


# Composing Complex Skills by Learning Transition Policies




Youngwoon Lee\* Shao-Hua Sun\* Sriram Somasundaram Edward S. Hu Joseph J. Lim

## Composing Complex Skills

- Primitive skills can be achieved by hard-coding, RL, and imitation learning
- We can compose a complex skill by sequentially executing primitive skills
- However, an ending state of a primitive skill may not be a good state to initiate the following skill



- A transition policy learns to smoothly connect primitive skills




## Proximity Reward

- A successful execution of a transition policy is defined by success of the following primitive skill
- The success/failure reward is too sparse to train a transition policy
- A proximity predictor learns to predict the proximity of a state to the initiation set

$$P(s) = \delta^{step(s)}$$


- We use the increase of predicted proximity to the initiation set as a dense reward

$$R_{proximity}(s_t, s_{t+1}) = P(s_{t+1}) - P(s_t)$$




The numbers inside the circles (states) represent transition steps to the initiation set

## Modular Framework with Transition Policies



## Training Transition Policies




- Train proximity predictors:  $L_P(\omega, \mathcal{B}^S, \mathcal{B}^F) = \frac{1}{2} \mathbb{E}_{(s, v) \sim \mathcal{B}^S} [(P_\omega(s) - v)^2] + \frac{1}{2} \mathbb{E}_{s \sim \mathcal{B}^F} [P_\omega(s)^2]$
- Train transition policies with proximity reward:  $R_{proximity}(\phi) = \mathbb{E}_{(s_0, s_1, \dots, s_T) \sim \pi_\phi} [\gamma^T P_\omega(s_T) + \sum_{t=0}^{T-1} \gamma^t (P_\omega(s_{t+1}) - P_\omega(s_t))]$


## Results

Baselines:




- (1) TRPO, PPO: train a policy with a hand-engineered reward for 5x longer time
- (2) TP-Task: train transition policies with a subtask completion reward
- (3) TP-Sparse: train transition policies with a sparse proximity reward
- (4) TP-Dense (ours): train transition policies with a dense proximity reward




Repetitive picking up




Patrol (Forward-Balance-Backward)



Repetitive catching




Hurdle (Forward-Jump)



Serve (Toss-Hit)



Obstacle course (Forward-Jump-Crawl)



Transition trajectories of "Repetitive picking up" and "Patrol"



## Code is available

- Code and videos are available at <https://youngwoon.github.io/transition>
- Corresponding author: Youngwoon Lee ([lee504@usc.edu](mailto:lee504@usc.edu))